

## **Operating Systems** **CMPSCI 377** ***Garbage Collection***

**Emery Berger**

University of Massachusetts Amherst



# Questions To Come

- **Terms:**
  - **Tracing**
  - **Copying**
  - **Conservative**
  - **Parallel GC**
  - **Concurrent GC**
- **Runtime & space costs**
  - **Live objects**
  - **Reachable objects**
  - **References**



# Explicit Memory Management

- **malloc/new**
  - allocates space for an object
- **free/delete**
  - returns memory to system
- Simple, but tricky to get right
  - Forget to **free**       $\Rightarrow$  memory leak
  - **free** too soon       $\Rightarrow$  “dangling pointer”
  - Double **free**, invalid **free**...



# Dangling Pointers

```
Node x = new Node ("happy");  
Node ptr = x;  
delete x;           // But I'm not dead yet!  
Node y = new Node ("sad");  
cout << ptr->data << endl;    // sad ☹️
```

- Insidious, hard-to-track down bugs



# Solution: Garbage Collection

- **Garbage collector** periodically scans objects on heap
  - No need to free
  - **Automatic memory management**
- Reclaims ***non-reachable*** objects
  - Won't reclaim objects until they're **dead** (actually somewhat later)



# No More Dangling Pointers

```
Node x = new Node ("happy");  
Node ptr = x;  
// x still live (reachable through ptr)  
Node y = new Node ("sad");  
cout << ptr->data << endl; // happy! 😊
```

So why not use GC  
**all the time?**



# Because This Guy Says Not To

**Linus  
Torvalds**

**GC sucks donkey brains  
through a straw from a  
performance standpoint.**

**to manage your  
memory.**



# Slightly More Technically...

- “GC impairs performance”
  - Extra processing
    - *collection, copying*
  - Degrades cache performance (*ibid*)
  - Degrades page locality (*ibid*)
  - Increases memory needs
    - *delayed reclamation*





# On the other hand...

- No, “GC enhances performance!”
  - Faster allocation
    - *pointer-bumping vs. freelists*
  - Improves cache performance
    - *no need for headers*
  - Better locality
    - *can reduce fragmentation, compact data structures according to use*



# Outline

- Classical GC algorithms
- Quantifying GC performance
  - A hard problem
- Oracular memory management
- GC vs. **malloc/free** bakeoff



# Classical Algorithms

- Three classical algorithms
  - **Mark-sweep**
  - **Reference counting**
  - **Semispace**
- Tweaks
  - **Generational garbage collection**
- Out of scope
  - **Parallel** – perform GC in parallel
  - **Concurrent** – run GC at same time as app
  - **Real-time** – ensure bounded pause times

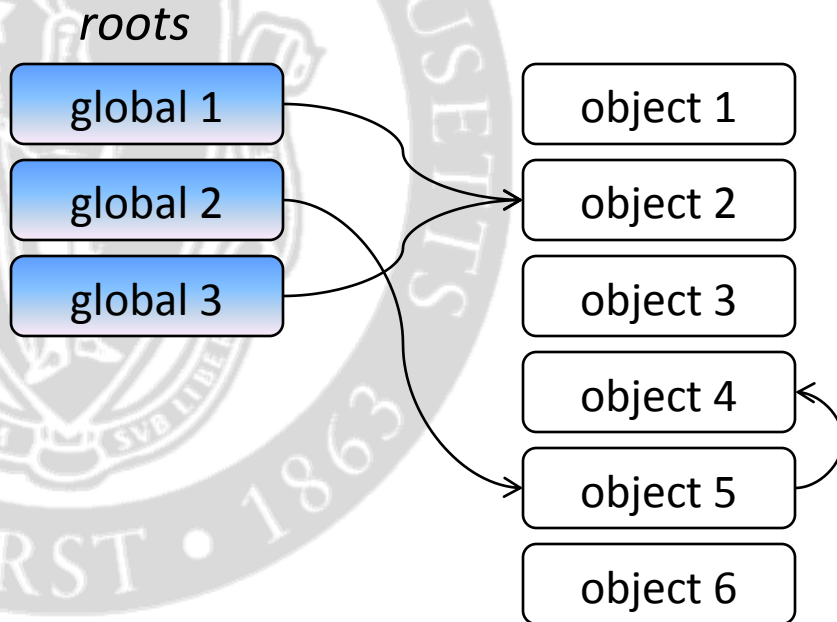


# Mark-Sweep

- Start with **roots**
  - Global variables, variables on stack & in registers
- Recursively visit every object through pointers
  - **Mark** every object we find (set **mark bit**)
- **Everything not marked = garbage**
  - Can then **sweep** heap for unmarked objects and free them



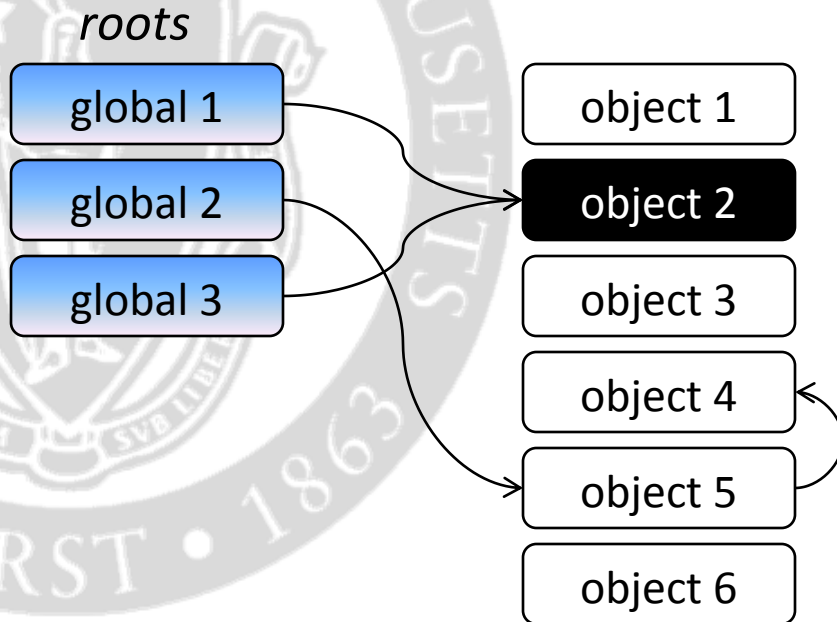
# Mark-Sweep Example



- Initially, all objects **white** (garbage)



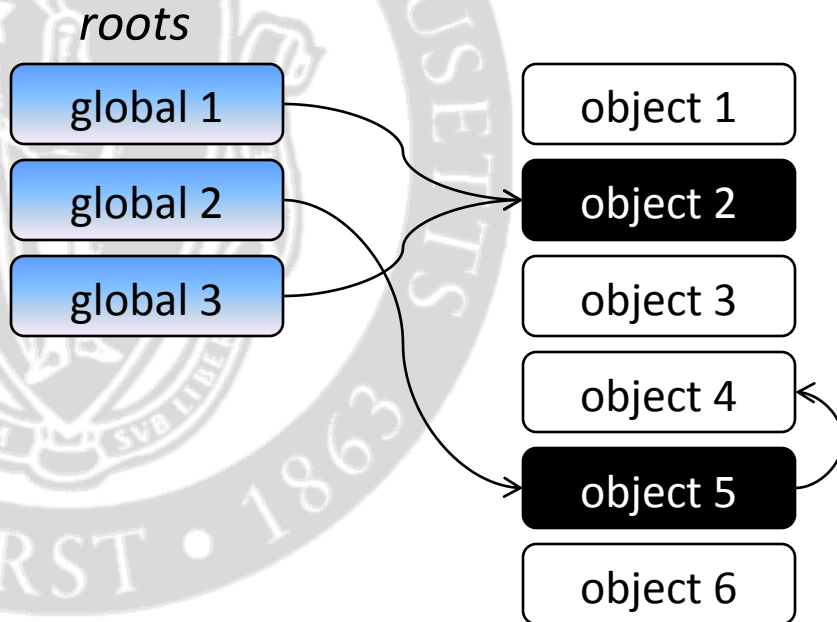
# Mark-Sweep Example



- Initially, all objects **white** (garbage)
- Visit objects, following **object graph**



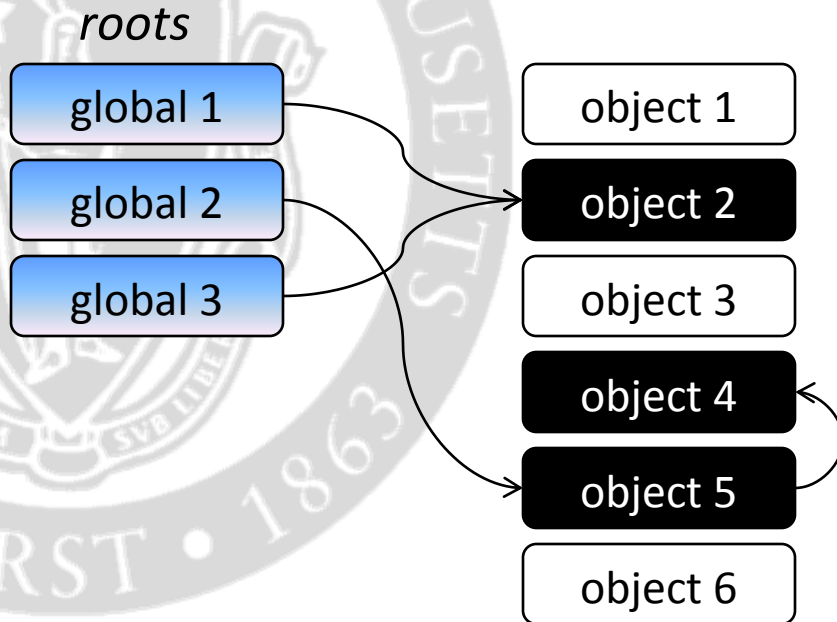
# Mark-Sweep Example



- Initially, all objects **white** (garbage)
- Visit objects, following **object graph**



# Mark-Sweep Example

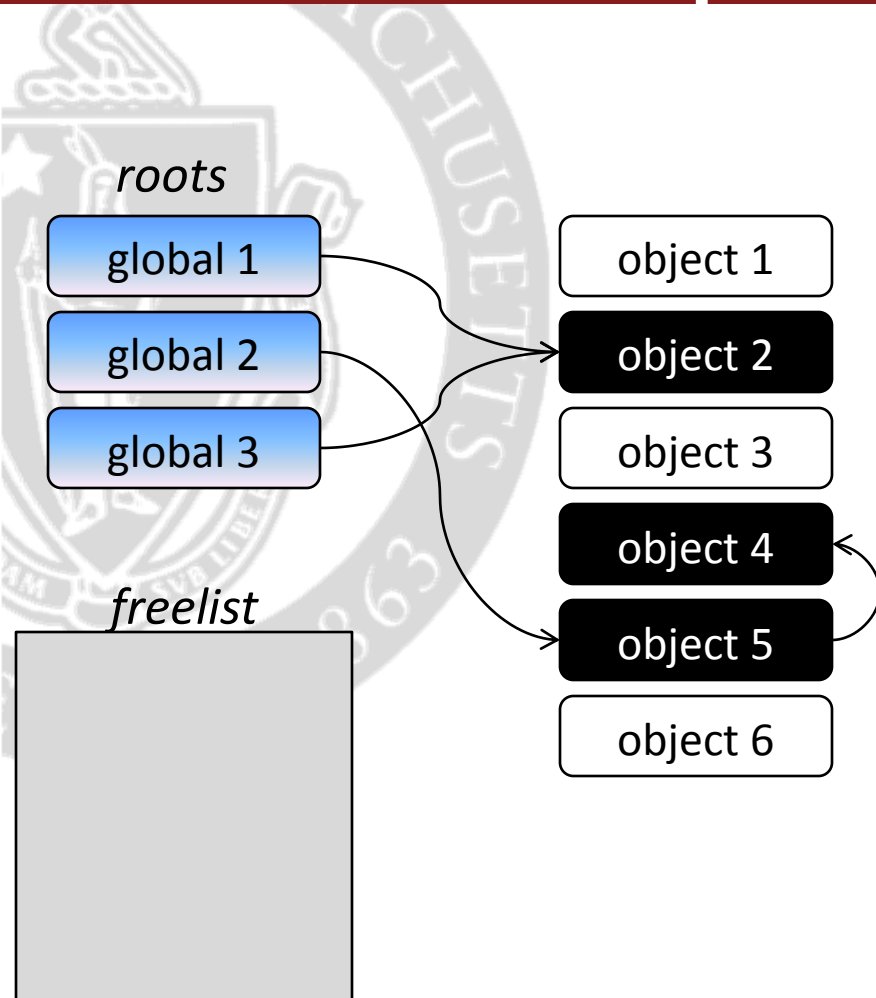


- Initially, all objects **white** (garbage)
- Visit objects, following **object graph**





# Mark-Sweep Example



- Initially, all objects **white** (garbage)
- Visit objects, following **object graph**
- Can sweep immediately or **lazily**

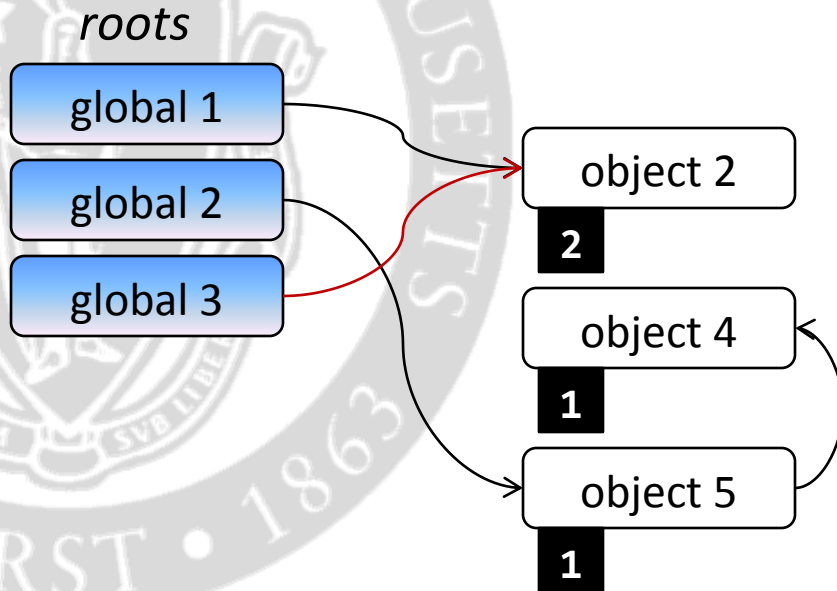


# Reference Counting

- For every object, maintain **reference count**  
= number of incoming pointers
  - $a \rightarrow ptr = x \Rightarrow \text{refcount}(x)++$
  - $a \rightarrow ptr = y \Rightarrow \text{refcount}(x)--;$   
 $\text{refcount}(y)++$
- Reference count = 0  $\Rightarrow$   
no more incoming pointers: **garbage**

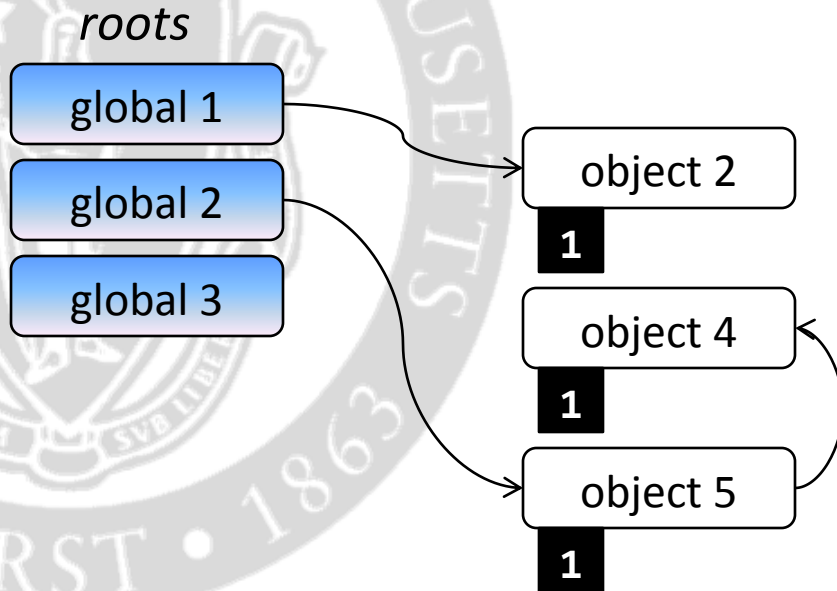


# Reference Counting Example



- New objects:  
ref count = 1

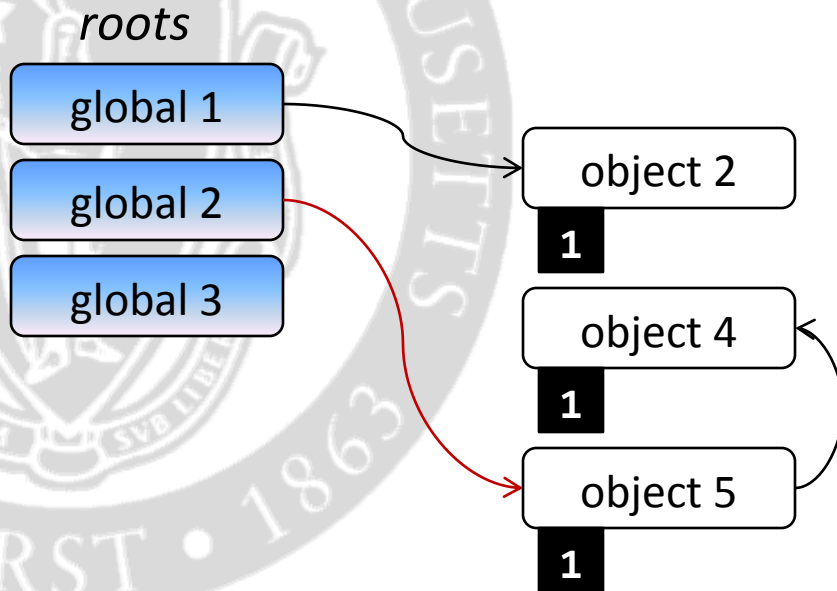
# Reference Counting Example



- New objects:  
ref count = 1
- Delete pointer:  
refcount--



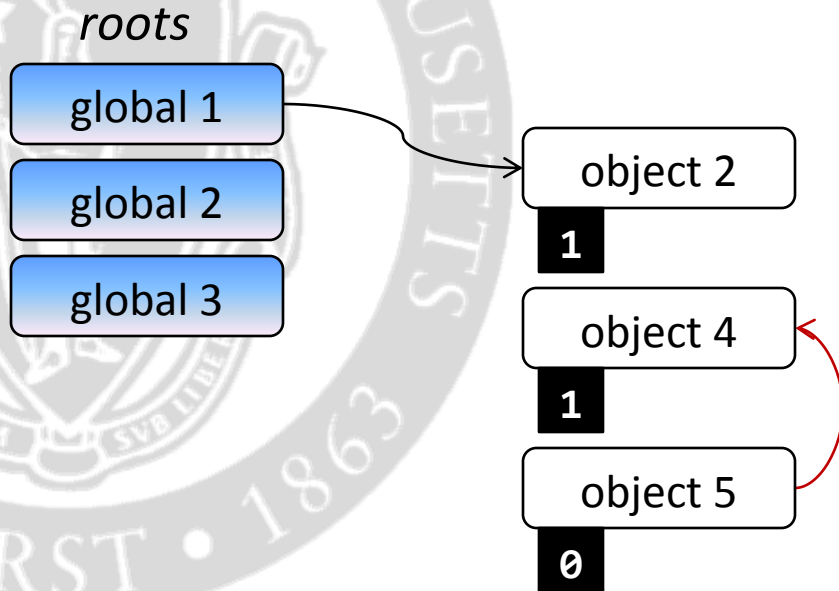
# Reference Counting Example



- New objects:  
ref count = 1
- Delete pointer:  
refcount--



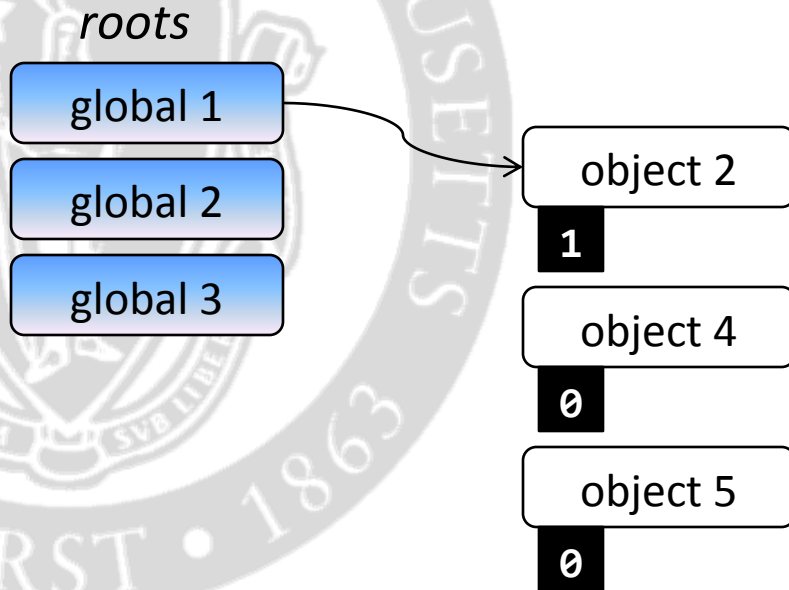
# Reference Counting Example



- New objects:  
ref count = 1
- Delete pointer:  
refcount--
  - And recursively  
delete pointers



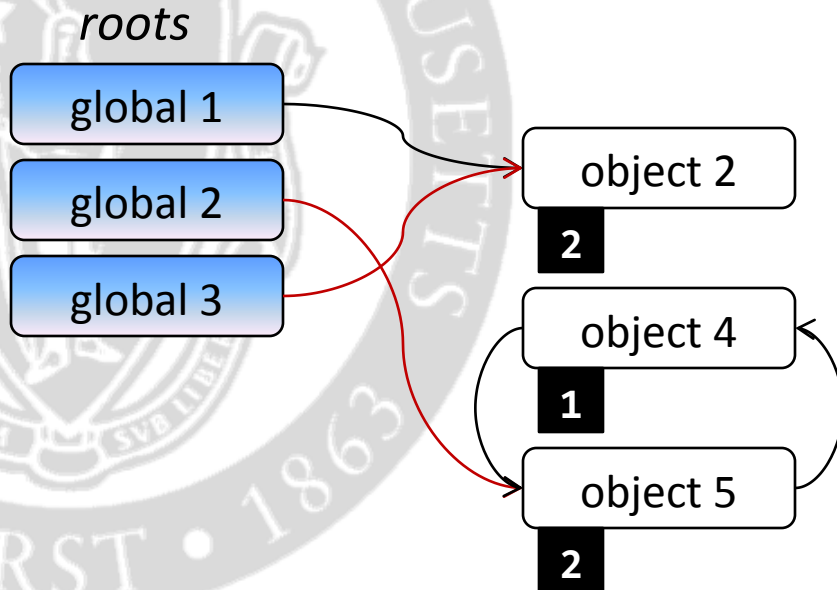
# Reference Counting Example



- New objects:  
ref count = 1
- Delete pointer:  
refcount--
  - And recursively  
delete pointers
- refcount == 0:  
put on freelist



# Cycles & Reference Counting

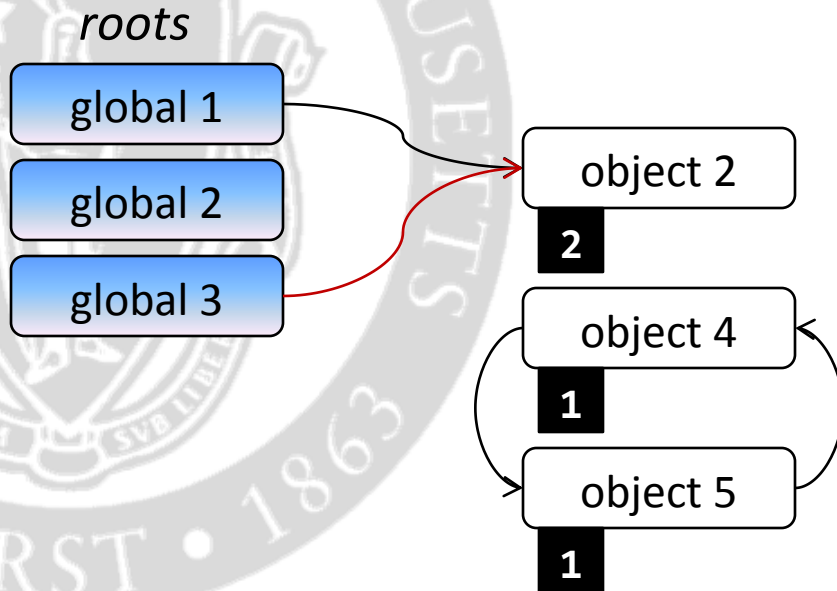


- **Big problem: cycles**






# Reference Counting Example



- **Big problem: cycles**



# Reference Counting Example



- **Big problem: cycles**
  - Cycles lead to unreclaimable garbage
- Need to do periodic tracing collection (e.g., mark-sweep)



# Semispace

- Divide heap in two **semispaces**:
  - Allocate objects from **from-space**
- When from-space fills,
  - **Scan** from roots through live objects
  - **Copy** them into **to-space**
- When done, *switch the spaces*
  - Allocate from leftover part of to-space



# Semispace Example

*from-space*



*to-space*



- Allocate in from-space
  - **Pointer bumping**




# Semispace Example

*from-space*



*to-space*



- Allocate in from-space
  - **Pointer bumping**



# Semispace Example

*from-space*




*to-space*



- Allocate in from-space
  - **Pointer bumping**





# Semispace Example



- Allocate in from-space
  - **Pointer bumping**
- Copy live objects into to-space



# Semispace Example




- Allocate in from-space
  - **Pointer bumping**
- Copy live objects into to-space
  - Leaves **forwarding pointer**






# Semispace Example



- Allocate in from-space
  - **Pointer bumping**
- Copy live objects into to-space
  - Leaves **forwarding pointer**




# Semispace Example



- Allocate in from-space
  - **Pointer bumping**
- Copy live objects into to-space
  - Leaves **forwarding pointer**



# Semispace Example



- Allocate in from-space
  - **Pointer bumping**
- Copy live objects into to-space
  - Leaves **forwarding pointer**
- **Flip** spaces; allocate from end



# Generational GC

- Optimization for copying collectors
  - **Generational hypothesis:**  
“most objects die young”
  - **Common-case optimization**
- Allocate into **nursery**
  - Small region
  - Collect frequently
    - Copy out **survivors**
- **Key idea:** keep track of pointers from **mature space** into **nursery**



# Generational GC Example

*mature space*



*nursery*



# Generational GC Example

*mature space*



*nursery*




# Generational GC Example


*mature space*



*nursery*




# Generational GC Example



- Copy out survivors (via roots & mature space pointers)




# Generational GC Example




- Copy out survivors (via roots & mature space pointers)

# Generational GC Example



- Copy out survivors (via roots & mature space pointers)

# Generational GC Example



- Copy out survivors (via roots & mature space pointers)
- Reset allocation pointer & continue

# Conservative GC

- **Non-copying** collectors for C & C++
  - Must identify pointers
    - “Duck test”: if it looks like a pointer, it’s a pointer
  - Trace through “pointers”, marking everything
- Can link with Boehm-Demers-Weiser library (“libgc”)



## GC vs. malloc/free



# Comparing Memory Managers

```
Node v = malloc(sizeof(Node));
v->data=malloc(sizeof(NodeData));
memcpy(v->data, old->data,
        sizeof(NodeData));
free(old->data);
v->next = old->next;
v->next->prev = v;
v->prev = old->prev;
v->prev->next = v;
free(old);
```

BDW  
Collector

Using GC in C/C++ is easy:



# Comparing Memory Managers

```
Node v = malloc(sizeof(Node));  
v->data=malloc(sizeof(NodeData))  
memcpy(v->data, old->data,  
        sizeof(NodeData));  
free(old->data);  
v->next = old->next;  
v->next->prev = v;  
v->prev = old->prev;  
v->prev->next = v;  
free(old);
```

BDW  
Collector

...slide in BDW and ignore calls to **free**.



# What About Other Garbage Collectors?

- Compares `malloc` to GC, but only *conservative, non-copying* collectors
  - Can't reduce fragmentation, reorder objects, etc.
- But: faster *precise, copying* collectors
  - Incompatible with C/C++
  - Standard for Java...





# Comparing Memory Managers

```
Node node = new Node();  
node.data = new NodeData();  
useNode(node);  
node = null;  
...  
node = new Node();  
...  
node.data = new NodeData();  
...
```

Lea  
Allocator

Adding **malloc**/**free** to Java:  
not so easy...



# Comparing Memory Managers

```
Node node = new Node();  
node.data = new NodeData();  
useNode(node);  
node = null;
```

`free(node)`  
?

```
...  
node = new Node();  
...  
node.data = new NodeData();  
...
```

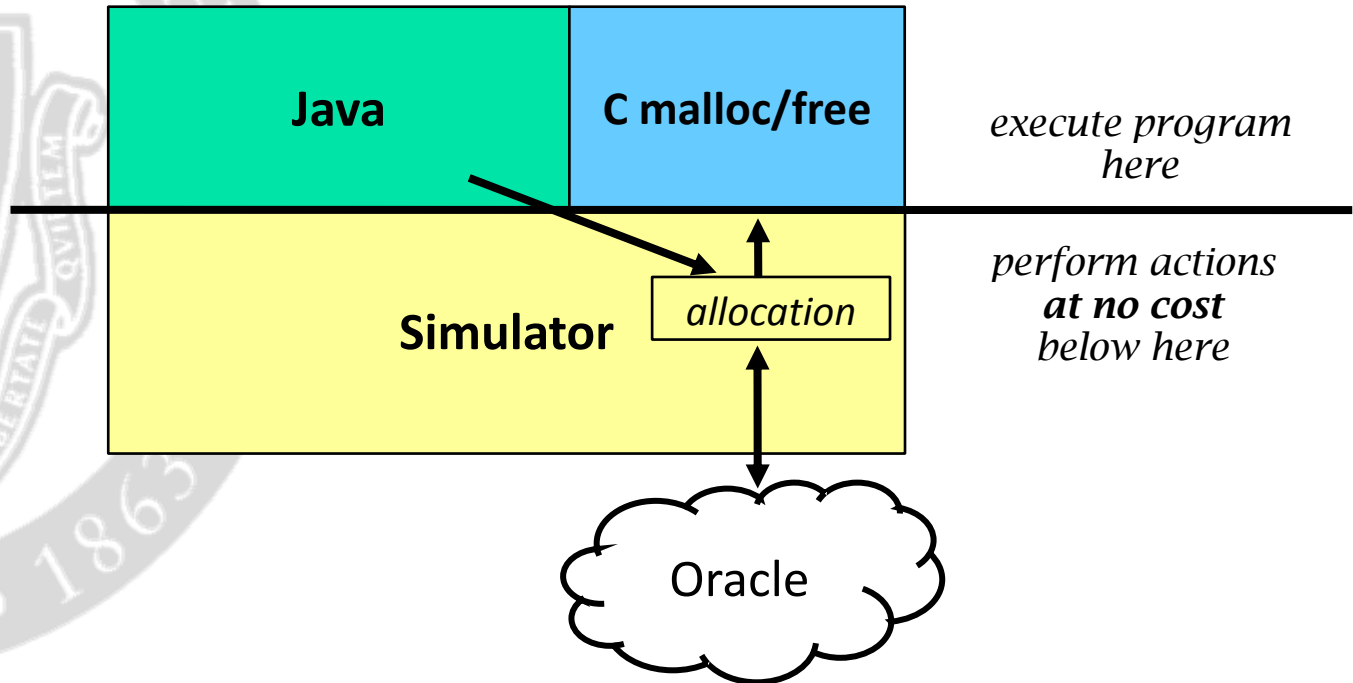
`free(node.data) ?`

Lea  
Allocator

... need to insert **free**s, but where?



# Oracular Memory Manager




- Consult oracle at each allocation
  - Oracle does not disrupt hardware state
  - Simulator invokes **free()** ...



# Object Lifetime & Oracle Placement

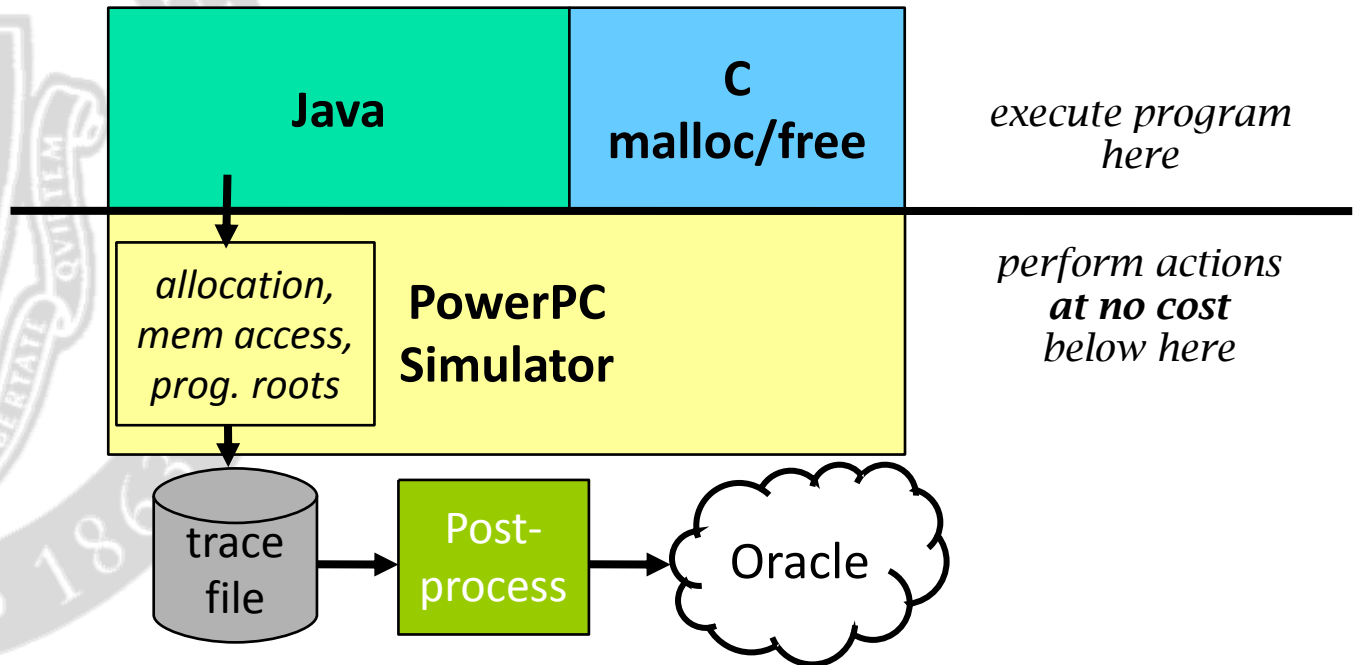
```
obj =  
new Object;
```



- Oracles bracket placement of **free**s
  - **Lifetime-based**: most aggressive
  - **Reachability-based**: most conservative



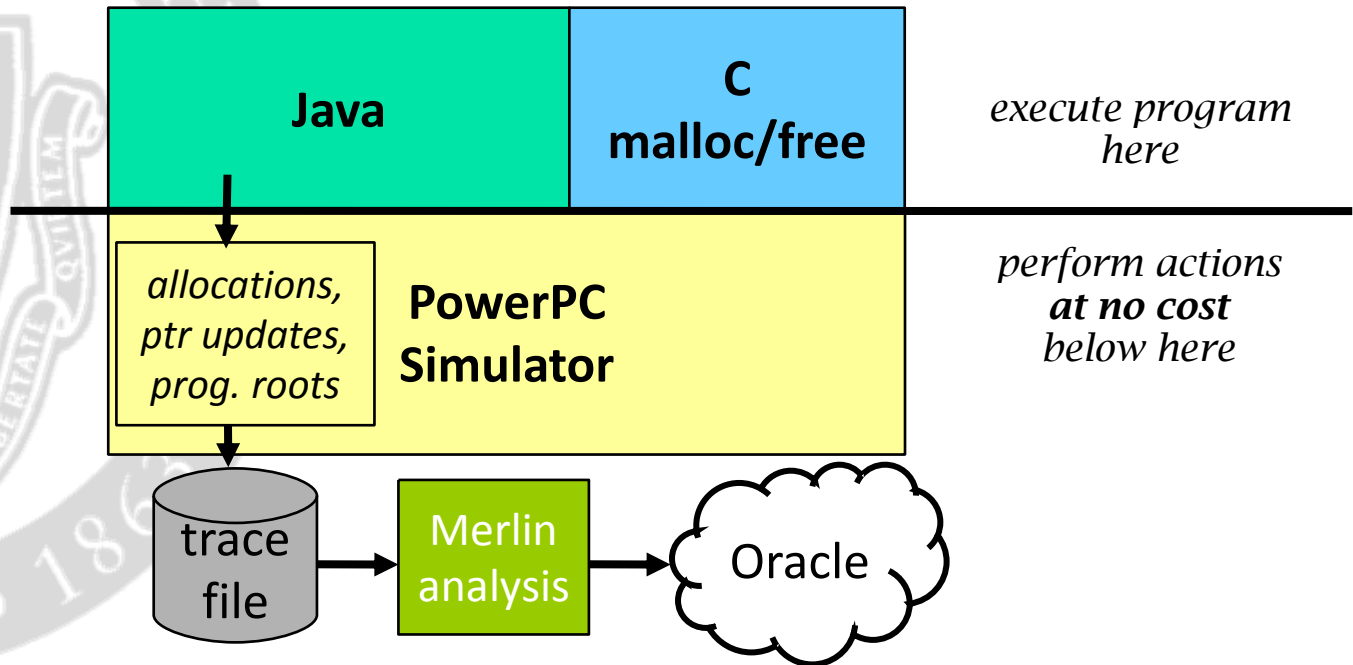
# Liveness Oracle Generation



- **Liveness:** record allocs, memory accesses



# Reachability Oracle Generation

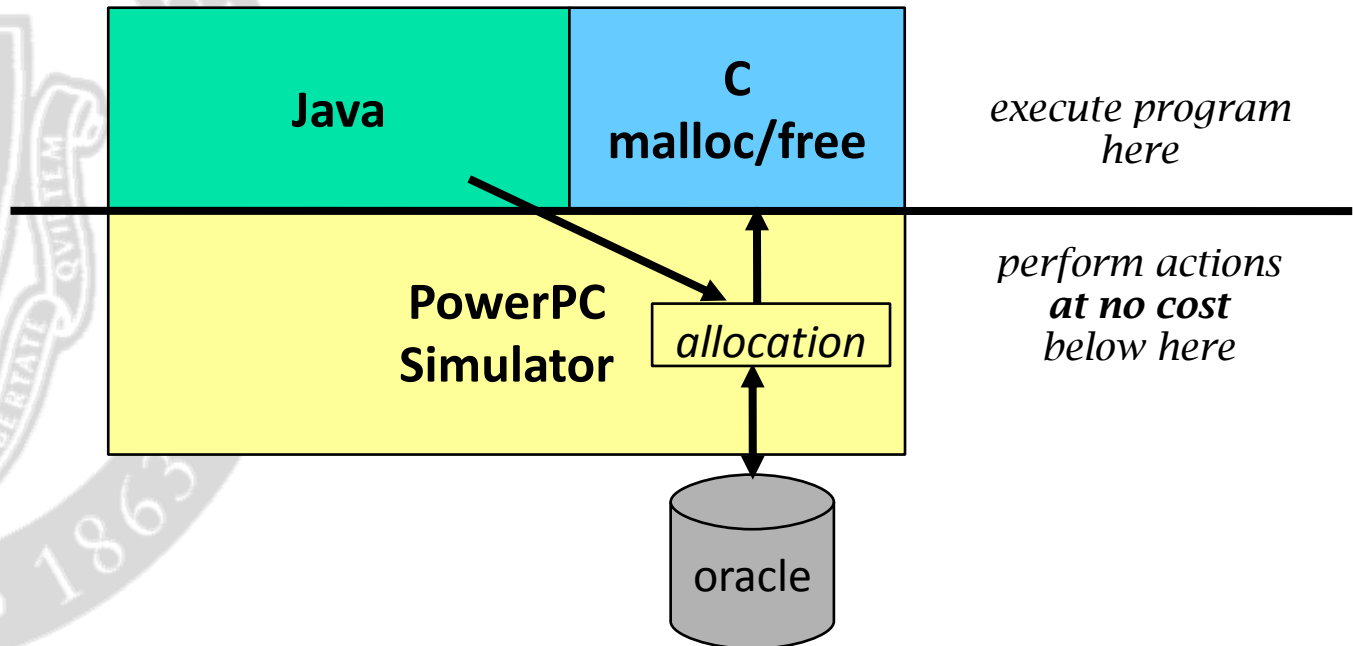


## ■ Reachability:

- Illegal instructions mark heap events (especially pointer assignments)



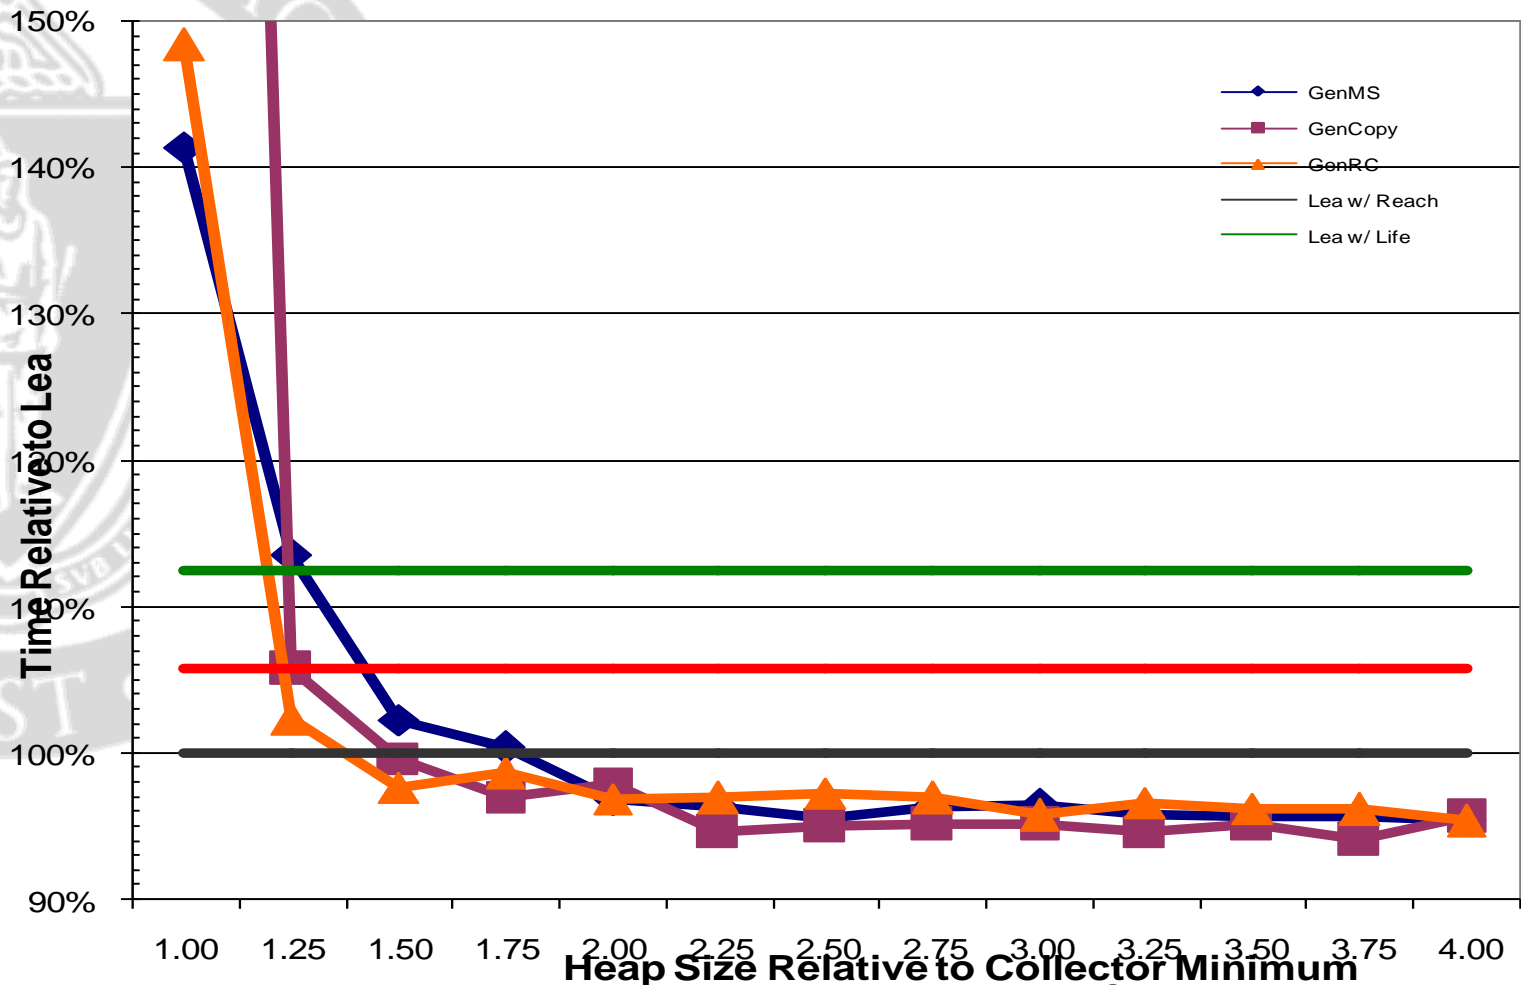
# Oracular Memory Manager



- Run & consult oracle before each allocation
  - When needed, modify instruction to call **free**
  - Extra costs (oracle access) hidden by simulator



# Execution Time for pseudoJBB

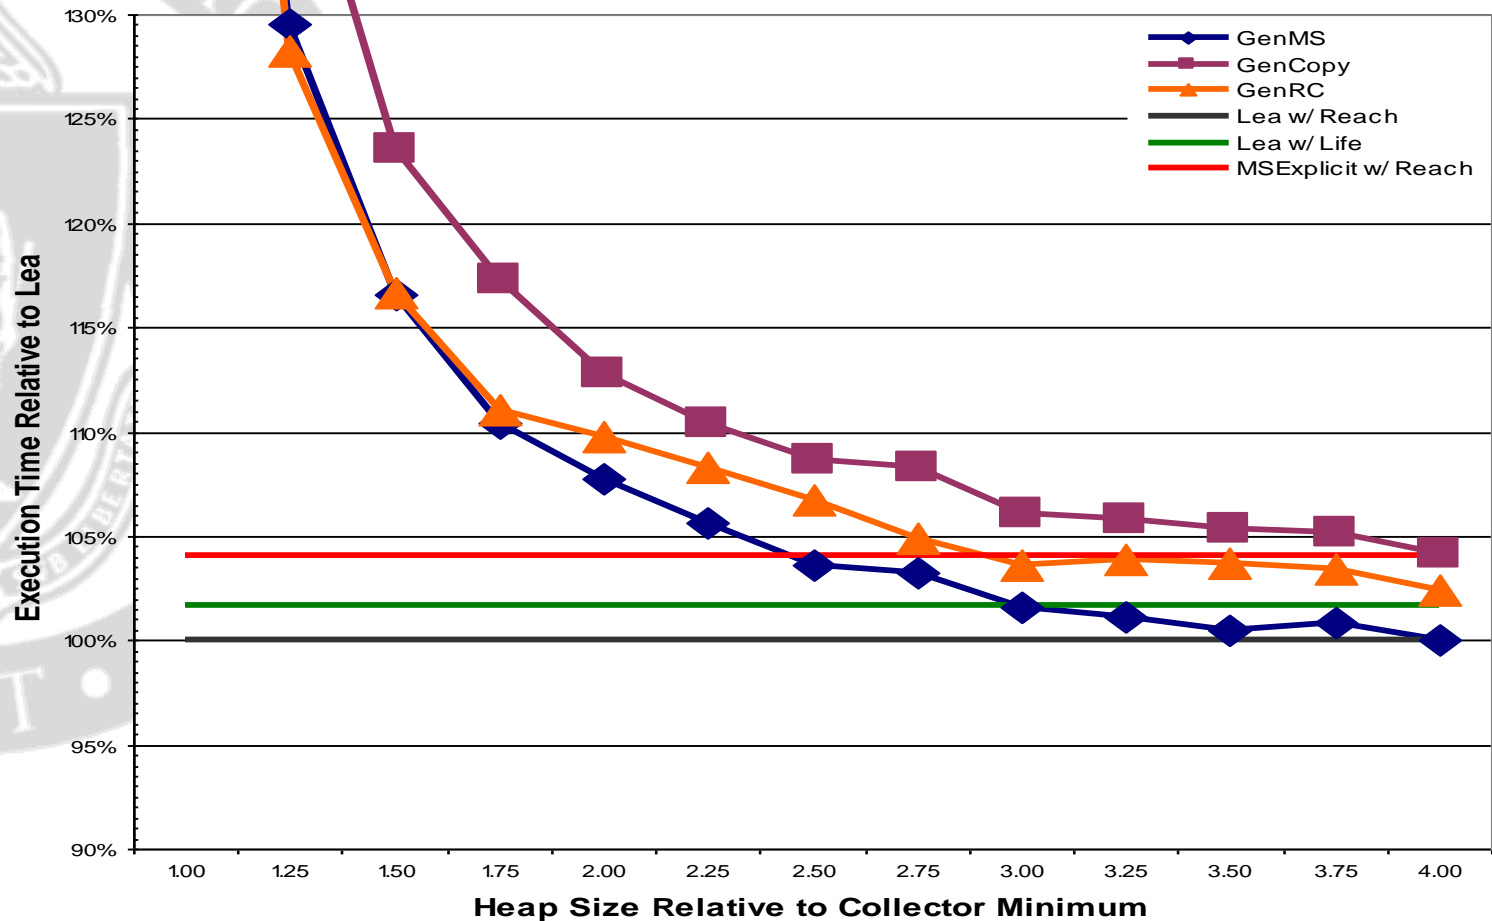


GC can be **faster** than **malloc/free**






# Geo. Mean of Execution Time



Trades space for time




# Footprint at Quickest Run



GC uses much more memory for speed




# Footprint at Quickest Run



GC uses much more memory for speed



# Javac Paging Performance



GC: poor paging performance

# Summary of Results

- Best collector equals Lea's performance...
  - Up to 10% faster on some benchmarks
- ... but uses more memory
  - Quickest runs require 5x or more memory
  - GenMS at least doubles mean footprint



# When to Use Garbage Collection

- Garbage collection fine if
  - system has more than 3x needed RAM
  - and no competition with other processes
  - **or** avoiding bugs / security more important
- Not so good:
  - Limited RAM
  - Competition for physical memory
  - Depends on RAM for performance
    - In-memory database
    - Search engines, etc.



# The End

